
MODULE *TwoPhase*

EXTENDS *TLC*

This specification is discussed in “Two-Phase *Commit*”, Lecture 6 of the TLA+ Video Course. It describes the Two-Phase *Commit* protocol, in which a transaction manager (*TM*) coordinates the resource managers (*RM*s) to implement the Transaction *Commit* specification of module *TCommit*. In this specification, *RMs* spontaneously issue *Prepared* messages. We ignore the Prepare messages that the *TM* can send to the *RMs*.

For simplicity, we also eliminate *Abort* messages sent by an *RM* when it decides to abort. Such a message would cause the *TM* to abort the transaction, an event represented here by the *TM* spontaneously deciding to abort.

CONSTANT RM	The set of resource managers
---------------	------------------------------

VARIABLES

<i>rmState</i> ,	<i>rmState</i> [<i>r</i>] is the state of resource manager <i>r</i> .
<i>tmState</i> ,	The state of the transaction manager.
<i>tmPrepared</i> ,	The set of <i>RM</i> s from which the <i>TM</i> has received “Prepared” messages.

msgs

In the protocol, processes communicate with one another by sending messages. For simplicity, we represent message passing with the variable *msgs* whose value is the set of all messages that have been sent. A message is sent by adding it to the set *msgs*. An action that, in an implementation, would be enabled by the receipt of a certain message is here enabled by the presence of that message in *msgs*. For simplicity, messages are never removed from *msgs*. This allows a single message to be received by multiple receivers. Receipt of the same message twice is therefore allowed; but in this particular protocol, that's not a problem.

$$Messages \stackrel{\Delta}{=} \quad$$

The set of all possible messages. Messages of type “Prepared” are sent from the *RM* indicated by the message’s *rm* field to the *TM*. Messages of type “Commit” and “Abort” are broadcast by the *TM*, to be received by all *RMs*. The set *msgs* contains just a single copy of such a message.

$$[type : \{\text{"Prepared"}\}, rm : RM] \cup [type : \{\text{"Commit"}, \text{"Abort"}\}]$$
$$TPT_{typeOK} \stackrel{\Delta}{=} \text{---}$$

The type-correctness invariant

$$\wedge rmState \in [RM \rightarrow \{\text{"working"}, \text{"prepared"}, \text{"committed"}, \text{"aborted"}\}]$$
$$\wedge tmState \in \{\text{"init"}, \text{"done"}\}$$
$$\wedge tmPrepared \subseteq RM$$
$$\wedge msgs \subseteq Messages$$
$$TPInit \triangleq$$

The initial predicate.

$$\wedge rmState = [r \in RM \mapsto \text{“working”}]$$
$$\wedge tmState = \text{"init"}$$
$$\wedge tmPrepared = \{\}$$
$$\wedge msgs = \{\}$$

We now define the actions that may be performed by the processes, first the TM 's actions, then the RM 's actions.

$TM_{RcvPrepared}(r) \triangleq$

The TM receives a “Prepared” message from resource manager r . We could add the additional enabling condition $r \notin tmPrepared$, which disables the action if the TM has already received this message. But there is no need, because in that case the action has no effect; it leaves the state unchanged.

$\wedge tmState = \text{“init”}$
 $\wedge [type \mapsto \text{“Prepared”}, rm \mapsto r] \in msgs$
 $\wedge tmPrepared' = tmPrepared \cup \{r\}$
 $\wedge \text{UNCHANGED } \langle rmState, tmState, msgs \rangle$

$TM_{Commit} \triangleq$

The TM commits the transaction; enabled iff the TM is in its initial state and every RM has sent a “Prepared” message.

$\wedge tmState = \text{“init”}$
 $\wedge tmPrepared = RM$
 $\wedge tmState' = \text{“done”}$
 $\wedge msgs' = msgs \cup \{[type \mapsto \text{“Commit”}]\}$
 $\wedge \text{UNCHANGED } \langle rmState, tmPrepared \rangle$

$TM_{Abort} \triangleq$

The TM spontaneously aborts the transaction.

$\wedge tmState = \text{“init”}$
 $\wedge tmState' = \text{“done”}$
 $\wedge msgs' = msgs \cup \{[type \mapsto \text{“Abort”}]\}$
 $\wedge \text{UNCHANGED } \langle rmState, tmPrepared \rangle$

$RM_{Prepare}(r) \triangleq$

Resource manager r prepares.

$\wedge rmState[r] = \text{“working”}$
 $\wedge rmState' = [rmState \text{ EXCEPT } ![r] = \text{“prepared”}]$
 $\wedge msgs' = msgs \cup \{[type \mapsto \text{“Prepared”}, rm \mapsto r]\}$
 $\wedge \text{UNCHANGED } \langle tmState, tmPrepared \rangle$

$RM_{ChooseToAbort}(r) \triangleq$

Resource manager r spontaneously decides to abort. As noted above, r does not send any message in our simplified spec.

$\wedge rmState[r] = \text{“working”}$
 $\wedge rmState' = [rmState \text{ EXCEPT } ![r] = \text{“aborted”}]$
 $\wedge \text{UNCHANGED } \langle tmState, tmPrepared, msgs \rangle$

$RM_{RcvCommitMsg}(r) \triangleq$

Resource manager r is told by the TM to commit.

$\wedge [type \mapsto \text{“Commit”}] \in msgs$
 $\wedge rmState' = [rmState \text{ EXCEPT } ![r] = \text{“committed”}]$
 $\wedge \text{UNCHANGED } \langle tmState, tmPrepared, msgs \rangle$

$RM\text{RcvAbortMsg}(r) \triangleq$

Resource manager r is told by the TM to abort.

$\wedge [type \mapsto \text{"Abort"}] \in msgs$

$\wedge rmState' = [rmState \text{ EXCEPT } ![r] = \text{"aborted"}]$

$\wedge \text{UNCHANGED } \langle tmState, tmPrepared, msgs \rangle$

$TPNext \triangleq$

$\vee TMCommit \vee TMAbort$

$\vee \exists r \in RM :$

$TM\text{RcvPrepared}(r) \vee RM\text{Prepare}(r) \vee RM\text{ChooseToAbort}(r)$

$\vee RM\text{RcvCommitMsg}(r) \vee RM\text{RcvAbortMsg}(r)$

The material below this point is not discussed in Video Lecture 6. It will be explained in Video Lecture 8.

$TPSpec \triangleq TPInit \wedge \Box[TPNext]_{\langle rmState, tmState, tmPrepared, msgs \rangle}$

The complete spec of the Two-Phase *Commit* protocol.

THEOREM $TPSpec \Rightarrow \Box TPTYPEOK$

This theorem asserts that the type-correctness predicate $TPTYPEOK$ is an invariant of the specification.

We now assert that the Two-Phase *Commit* protocol implements the Transaction *Commit* protocol of module $TCommit$. The following statement imports all the definitions from module $TCommit$ into the current module.

INSTANCE $TCommit$

THEOREM $TPSpec \Rightarrow TCSpec$

This theorem asserts that the specification $TPSpec$ of the Two-Phase Commit protocol implements the specification $TCSpec$ of the Transaction *Commit* protocol.

The two theorems in this module have been checked with TLC for six RM s, a configuration with 50816 reachable states, in a little over a minute on a 1 GHz PC .
